

Copilot Credits

The September 1 Exposure

Messaging, communications and field enablement response to the usage based billing transition.

61%

Drop in Business allowance

64%

Drop in Enterprise allowance

DECISION REQUIRED

1 September 2026

TO

Copilot GTM leadership

FROM

Naveen Raj Kanagaraj, Product Marketing

DATED

18 August 2026

Internal recommendation

Diagnosis, repositioning, communications sequence and field enablement for the shift to usage based billing.

*

2026

The transition landed. The invoice has not.

1 JUNE

Premium requests replaced by metered AI Credits across every plan.

PRICE

Base plans unchanged. Business \$19, Enterprise \$39 per user, per month.

SCOPE

Completions stay free. Chat, CLI, agents and code review are metered.

COVER

Promotional credits for existing customers run June, July and August only.

SIGNAL

The backlash is almost entirely individuals. Enterprise is still silent.

We did not raise prices. We removed the ability to forecast them.

A COST INCREASE

is absorbed.

Negotiated at renewal. It stays inside engineering.

A FORECASTING FAILURE

is escalated.

It reaches finance, triggers procurement review and a spend freeze.

**We produced the second and are being criticised for the first.
Every fix that adjusts price solves the wrong problem.**

Credits burn on tokens priced at each model's API rate. Frontier and lightweight models differ by an order of magnitude for identical work, and the person choosing is a developer mid task with no budget context.

The safety net came out in the same release as the meter.

BEFORE

Allowance runs out.
Model falls back.
Work continues, degraded.

Graceful. No decision required.

AFTER


Allowance runs out.
No fallback exists.
Uncapped spend, or work stops.

Binary. Decided in advance, by someone else.

THE CONSEQUENCE WE HAVE NOT MESSAGED

The field is selling budget caps as the answer to unpredictability. The cap is also an off switch. We tell admins that predictable spend is one setting away, without telling them that setting can stop their developers mid task.

On 1 September, included usage falls by two thirds.



BUSINESS

61% less

\$49 to \$19 per user, per month


ENTERPRISE

64% less

\$109 to \$39 per user, per month

Promotional credits are additive. The \$30 and \$70 sit on top of the base allowance, not in place of it.

What 1,000 Enterprise seats pay in September.



WORST CASE

\$70K

of unplanned monthly overage, landing on an October invoice nobody forecast.

READ THIS AS A RANGE

September consumption is unknown to us today. The point is not the estimate. It is that the range is wide, and the customer cannot see where inside it they will land.

Assumes September usage holds at the June to August rate, which promotional credits themselves inflated. Directional range, not a forecast. Workings in appendix.

From predictable seat to governed platform.

WAS

"One price per developer.
Stop thinking about it."

IS

"Spend you can see, attribute
and govern before it happens."

POSITIONING STATEMENT / PLATFORM ENGINEERING LEAD

For the platform lead who now owns a variable AI budget they are accountable for, Copilot is the only AI coding platform that meters agentic work at the cost centre rather than the seat, so spend is attributable, capped and defensible in a budget review.

What it deliberately does not claim: that the bill is predictable. It is not, and claiming otherwise breaks trust a second time.

One reposition. Three translations.

DEVELOPER

Sees the meter

"See the cost before you commit."

PROOF

Completions stay free and unmetered.
Credit cost is shown at the point you pick a model.

CHANNEL

In product, docs, changelog

PLATFORM LEAD

Owns allocation

"Allocate by cost centre, not by seat."

PROOF

Pooled usage removes stranded capacity.
Budgets set at enterprise, cost centre or user level.

CHANNEL

Admin console, CSM, webinar

FINANCE

Sees the variance

"You choose the ceiling in advance."

PROOF

On exhaustion the organisation decides:
continue at published rates, or cap.

CHANNEL

Procurement FAQ, renewal brief

Everyone has this problem. We are first to show the customer.

Agentic coding consumes frontier inference. That cost is real for every vendor. The only question is who absorbs the variance, and for how long.

FLAT RATE SEAT

Most of the category today

Predictable invoice. Easy to buy.

The vendor absorbs variance, then rations quietly through throttles, fair use limits and downgrades at peak. Spend cannot be attributed to a team, because there is no spend to attribute.

METERED AND GOVERNED

Copilot, from 1 June 2026

Spend is attributable and capped in advance.

The invoice moves. The trust cost is paid up front and in public, which is where we are standing now.

Speculation: flat rate agentic pricing does not survive current inference economics. If competitors meter within 12 to 18 months, their customers arrive with no allocation model at all.

Land the cliff before it lands on them.

19 AUG	Admin notification	Every admin receives their own June to August consumption beside their standard September allowance.
19 AUG	Field briefed first	AEs and CSMs get the cap trade off before customers get the email.
25 AUG	Guided configuration	CSM led budget setup, ranked by projected gap rather than by ARR.
1 SEP	Say it plainly, in public	Changelog and docs state that promotional credits have ended. Silence converts a billing change into a trust story.
EARLY OCT	First unsubsidised invoice	Outreach to every account above forecast before they open it.

The objection that decides Q4 renewals.

"We cannot budget this. The variance is too wide."

WHAT WE SAY NOW

"Set a cap. Then you will never see an unplanned overage."

True, and incomplete. It omits what happens at the cap. An admin who sets one on our advice and finds their team blocked in week three will not call that a configuration issue.

WHAT TO SAY INSTEAD

"You are choosing between a spend ceiling and a throughput ceiling."

- Size the cap against observed June to August usage, not the base allowance
- Pool across cost centres so one team's spike cannot block another
- Alert well below the ceiling so the admin decides, not the meter
- Run the projection sheet with them, using their own actuals, not a worked example

One ask. Marketing cannot close this alone.

Budget controls, pooled usage and the preview bill all shipped. This is not a request for tooling we lack.

THE ASK

Restore a degraded capability floor at the cap.

When the pool is exhausted and the organisation has chosen to cap, route to a low cost model at reduced concurrency instead of stopping. The admin keeps a hard spend ceiling. The developer keeps a working tool.

WHY THIS AND NOT A DISCOUNT

A taper postpones the conversation by one cycle and costs revenue in every cycle after. A floor is paid for once.

IF THIS IS DECLINED

The claim weakens from governed to attributable. Marketing's fallback is alerting well ahead of the ceiling and a named approval path, which manages the stop rather than preventing it.

What this costs us.

WE GIVE UP

The predictability message.

Competitors run flat rate against us for at least two quarters and it works. We should expect to lose price sensitive mid market deals, and should not pretend otherwise in the forecast.

WE KEEP

A claim that stays true.

Attribution and governance can be demonstrated in a console today and will not be retracted. The alternative fails on the first spiky month and costs the account, not just the deal.

THE BET

That in enterprise software, trust compounds and price advantages do not. We went first, we are absorbing the reaction in public, and the allocation layer becomes table stakes the moment the category follows. If that reading is wrong, this is the wrong strategy, and the next slide says so by November.

Four numbers, and what would prove me wrong.

01**Budget configuration rate**

Organisations with a budget set at any level, by 30 September.

02**Cap proximity in week two**

Share of organisations whose usage approaches their own ceiling within fourteen days. The only signal that arrives early enough to recalibrate before the invoice.

03**Forecast gap at first invoice**

Admin projected against actual October spend. The direct measure, and the number I would be held to.

04**Billing cited in Q4 churn**

Non renewals naming predictability. Lagging and contaminated. Read alongside 03, never alone.

Falsification: If the October gap is wide and configuration is high, the diagnosis was wrong and this memo should be withdrawn, not iterated.

Tell them before the invoice does.

Reposition from predictable seat to governed platform, name the cap trade off before customers discover it, and ask Product for a capability floor so governance does not mean stoppage. The promotional window ends in two weeks either way.

01

Say it publicly on 1 September

02

Brief the field before the customer

03

Ship a floor at the cap

Decision required by 1 September 2026.

Working, sources and assets.

A	Assumptions and sources	Every figure traced to the GitHub announcement, with derivations.
B	Account risk segmentation	Ranking logic for outreach: projected gap, not ARR.
C	Full messaging matrix	Proof points, disqualifiers and phrases to avoid.
D	Admin notification email	Draft copy for the 19 August send.
E	Objection one pager	Cap trade off briefing for AEs and CSMs.
G	Spend projection sheet	CSM calculator: an account's own actuals against its September allowance.
F	Research basis and limits	What the evidence is and what it cannot support.

Assumptions and sources.

Business plan	\$19 per user per month, including \$19 in monthly AI Credits	GitHub
Enterprise plan	\$39 per user per month, including \$39 in monthly AI Credits	GitHub
Promotional credits	Additional \$30 and \$70 per user, June to August only	GitHub
Unmetered	Code completions and Next Edit suggestions, all plans	GitHub
Fallback	Discontinued. Usage governed by credits and budget controls	GitHub
Budget controls	Enterprise, cost centre and user level. Included usage pooled	GitHub
Sensitivity model	1,000 seats at (\$109 less \$39) across 50, 75 and 100% of usage	Derived
Not verified	September consumption, budget adoption, competitor pricing	Unknown